


Improvements to the ‘dvir’ Package

Alexander van der Voorn



Bachelor of Science (Honours)
Department of Statistics
The University of Auckland
New Zealand

Contents

1	Executive summary	3
2	Introduction	4
2.1	Where this project fits in	5
3	Background	6
3.1	\TeX	6
3.2	DVI	6
3.3	The (pre-existing) <code>dvir</code> package	6
4	Code speed (part 1) - removing redundant font sweeps	8
5	Code speed (part 2) - font caching	13
5.1	Profiling environment specifications	14
6	Linear gradient fills	16
6.1	<code>TikZ</code> and <code>dvir</code>	16
6.2	Implementing <code>TikZ</code> linear gradient fills in <code>dvir</code>	16
7	Text baselines	22
7.1	The problem	22
7.2	Implementation	22
7.3	Our potential solutions	23
7.4	Discussion of algorithms	29
8	Conclusion	30
9	References	31
10	Apendices	32

1 Executive summary

\TeX has many desirable features, such as math equation formatting and typesetting, which are currently not available in R. The `dvir` package extends R to draw \TeX output on R graphics. By extending the usability and functionality of the `dvir` package it better streamlines the workflow of making R graphics with some \TeX output - rather than making some graphics in R and exporting to \LaTeX itself or an image editor like Photoshop to add \TeX or \TeX -like output, this can all be done in R through a simple interface from the `dvir` package. Three improvements to the `dvir` package are investigated and discussed in this report.

Running functions from the `dvir` package took a long time which. The speed of the package was increased by making two changes - removing redundant code that meant the `dvir` package was searching for fonts three times instead once, and allowing font information to be cached so once font information was found the `dvir` package did not need to search again for that font.

The latest version of R adds functionality of linear and radial gradient fills. This means gradient fill information from \TikZ graphics, a \TeX package, could be generated by `dvir` and manipulated to be fed into R. This was not implemented as the complexity of the gradient definition from \TikZ resulted in “warped” gradient fills which are not possible in R.

The `dvir` package currently performs any alignment based on the bounding box of the text but aligning with the baseline of the text is more desirable. \TeX output does not explicitly state a baseline value for a piece of text so some algorithms for determining the text baselines were implemented in an R function to test their usefulness. One of these algorithms performs particular well for all tested examples, including multi-line text.

The magnitude of the speed increase alone makes `dvir` much more practical as it allows much faster testing of code and allows for more complicated graphics with many calls to the `dvir` package. The ability to align text to its baseline means graphics with the addition of \TeX , for example in lattice plot headers, can be aligned with one another and not require “touching up” outside of R.

This report, source code and associated files can be found on GitHub¹.

¹https://github.com/ajvandervoorn/honours_dissertation

2 Introduction

R has the ability to display mathematical symbols and equations in graphics using the “plotmath” feature, interpreting everything within a call to `expression()` as a mathematical equation.

```
mu <- 1:5
opar <- par(mar = par()$mar + c(0, 1, 0, 0))
plot(mu, mu ^ 2 / 2, xlab = expression(mu), ylab = "", yaxt = "n")
axis(2, las = 1)
mtext(expression(frac(mu ^ 2, 2)), side = 2, line = 3, las = 1)
```




Figure 1: A plot with axis labels made using ‘`expression()`’.

This provides us with most of the symbols used for equations, such as brackets and fractions, and formats them in a layout resembling \TeX , but it is limited in its fonts. Compare the y-axis label above with how it looks when created by \LaTeX in figure 2.

$$\frac{\mu^2}{2}$$

Figure 2: The y -axis label from above, when created in \LaTeX

The difference is stark and there are several approaches in R which can get us closer to the \LaTeX result Murrell (2018).

- **[Example:** use `extrafont` and `fontcm` packages, embed CM fonts in PDF]
- **[Example:** use `tikzDevice` package, which creates PGF/TikZ version of plot (and as such converts all text in plot to LaTeX (including labels))]

What we want is a middle ground - being able to harness the power of \TeX and its typesetting capabilities on our choice of text or equation in R graphics. This is where the `dvir` package comes in - providing a simple

user interface, in the style of the R `grid` graphics package [Reference R grid graphics here], by way of the `grid.latex()` function:

- **[Example:** example of grid-based plot, changing labels and/or title with `grid.latex()`. Maybe `ggplot2`?]

2.1 Where this project fits in

The `dvir` package already worked really well in a lot of cases. There were however plenty more desirable features of \TeX and its extensions though that had not yet been implemented by `dvir`. The power of this package is from ensuring it is comprehensive enough to meet a user’s entire \TeX needs in R graphics without having to leave R to do annotations in \LaTeX itself (or Photoshop/Illustrator!).

By keeping things “in R” users only need to learn R (and basic \TeX) code to create their graphics and their work is in one place and easily reproducible. It may not be realistic to *completely* replicate \TeX in R, however there were several aspects of the package identified as having a lot of potential to greatly increase its usefulness. The aspects identified were:

- the speed of the package - anecdotally it took a while to generate graphics, especially if there were many `grid.latex()` calls
- expand `dvir`’s capability of creating TikZ drawings by adding support for linear gradient fills
- adding the ability to align text from `grid.latex()` to a baseline - the natural line on which characters sit

3 Background

3.1 T_EX

T_EX is a program to format and typeset text, and includes some basic macros to do this. L^AT_EX is a higher-level implementation of T_EX, basically consisting of a lot more macros, creating a much more user-friendly interface to T_EX. For example, L^AT_EX allows one to create a document with numbered sections, title pages and bibliographies without having to write complicated T_EX macros themselves. There are other extensions to T_EX that do similar things to L^AT_EX too.

3.2 DVI

A T_EX or L^AT_EX file is just plain text, so there needs to be a step to translate this plain text to what you will see on a formatted document on a screen or page. A DVI (DeVice Independent) file is a binary file *describing* the layout of the document. For example, the height of the page, what characters to display and where, and the fonts to be used.

3.3 The (pre-existing) dvir package

In a simplified form, `dvir` works by providing a high level function, `grid.latex()`, to call with the T_EX code of the expression or text to be displayed.

```
library(dvir)
grid.latex("$x - \\mu$")
```

$$x - \mu$$

Figure 3: Using the ‘dvir’ function ‘grid.latex()’

The following steps are taken when `grid.latex()` runs:

1. A T_EX document is created with the expression and a changeable default preamble and postamble.
- `[**Example:** TeX document with pre- and post-ample]`
2. This TeX document is then processed using the local T_EX installation to create a DVI (DeVice Independent) file.
3. The DVI file is read into R. As DVI files are binary they are not easily readable by humans but the ‘dvir’ function ‘readDVI()’ translates the DVI file into readable text.
- `[**Example:** Extract of DVI file using ‘readDVI()’ (not the whole thing, just the bit relevant to our example $(x - \mu)$)]`
4. Three "sweeps" of the DVI file are completed to extract necessary information about what to display in R (and where and how to display it):

- Font sweep: Gather the names of all fonts used in the DVI file and locate the relevant font files on the local machine. The font information is stored in a R list as well as a ‘fontconfig’ file.
 - Metric sweep: To determine the overall bounding box (size) of the expression to display. This bounding box is used to create a ‘grid’ viewport which can encompass the entire $\text{T}_\text{E}\text{X}$ passed to ‘grid.latex()’ expression using the native DVI coordinates.
 - Grid sweep: Convert all text and symbols into *grobs* (grid graphical objects)
5. These grobs are then displayed in the R graphics device as per the ‘grid’ package.

4 Code speed (part 1) - removing redundant font sweeps

In the introduction of this report the case for the `dvir` package was motivated with a simple example of a mathematical equation. `dvir` can be used on a larger scale too.

```
grid.latex("\\dots", x = 0.44, y = -0.1, default.units="native")
grid.latex("$Y_* =$", x = 0.5, y = 1.1, default.units="native")
grid.latex("$a_1$", xpos[1], y = -0.1, default.units="native")
grid.latex("$a_2$", xpos[2], y = -0.1, default.units="native")
grid.latex("$a_{L_A}$", xpos[3], y = -0.1, default.units="native")
grid.latex("$Y_{\\pi} \\mid Y_{\\pi} \\notin \\mathcal{A}$", x = xpos[4], y = -0.1, default.units="native")
grid.latex("$\\omega_1$", x = 0.18, y = 0.50, default.units="native")
grid.latex("$\\omega_2$", x = 0.32, y = 0.50, default.units="native")
grid.latex("$\\dots$", x = 0.44, y = 0.50, default.units="native")
grid.latex("$\\omega_{L_A}$", x = 0.54, y = 0.50, default.units="native")
grid.latex("$1 - \\sum_{s=1}^{L_A} \\omega_s$", x = 0.95, y = 0.50, default.units="native")
```




Figure 4: Using nine ‘`grid.latex()`’ calls to create a graphic.

The example in figure ?? uses nine calls to `grid.latex()` and was created by a University of Auckland lecturer using the `dvir` package to help write an assignment.

One of the first things investigated in the package was the speed of running the code. Anecdotally, generating any R graphic with non-trivial \TeX , like that in figure 4, took a long time so it was desirable to see if we could speed it up.

To look into this the first task was to profile the existing code to let us see where in the package time was being spent. This was in `dvir` version 0.2-1.

We visualised the profiling results using `profvis::profvis()`.

We can see the function call stack in figure @ref(fig:profilingSimpleProfvis_0.2-1). At the bottom is the call to `grid.latex()`, which immediately calls `grid.draw()` which in turn calls `latexGrob()`. This calls `readDVI()` for about the first 20ms, then `dviGrob()` for the remaining time to the end of the original `grid.latex()` function call, and so on up the function call stack.




Figure 5: Screenshot of `profvis::profvis()` output for the code `grid.latex("$x - \mu$")` in `dvir` version 0.2-1.




Figure 6: Screenshot of `profvis::profvis()` output for the code creating figure 4 highlighting the time spent in `engine$fontDef`.

The `profvis::profvis()` output for our more complicated example, in figure @ref(fig:profilingYeeProfvis_0.2-1_highlight) reveals most of the time to create the figure is in `grid.latex()`. Note that the code to draw the arrows and the “(a)” in this example is so quick it occupies the very skinny call stack on the far left of the graph. `grid.latex()` and its subsequent function calls, on the other hand, take up most of the time required to produce the example.

In figure @ref(fig:profilingYeeProfvis_0.2-1_highlight) some blocks in the call stack have been highlighted - these are related to the `engine$fontDef` operation occurring. This is a part of the “font sweep”, as was described in the introduction to `dvir` in section 3.3.

In the top left corner of figure @ref(fig:profilingYeeProfvis_0.2-1_highlight) we are told the aggregate time spent with `engine$fontDef` is 1830ms. Compared to the total time of this run (a total of about 2700ms), `dvir` is spending a *lot* of time doing these font sweeps.

What was interesting though was that after the actual font sweep the following sweeps for the metric and grid information *also* called `engine$fontDef`. As the point of the font sweep is that it finds all the font information to be used later on the following metric and grid sweeps should not need to “re-sweep” for the fonts.

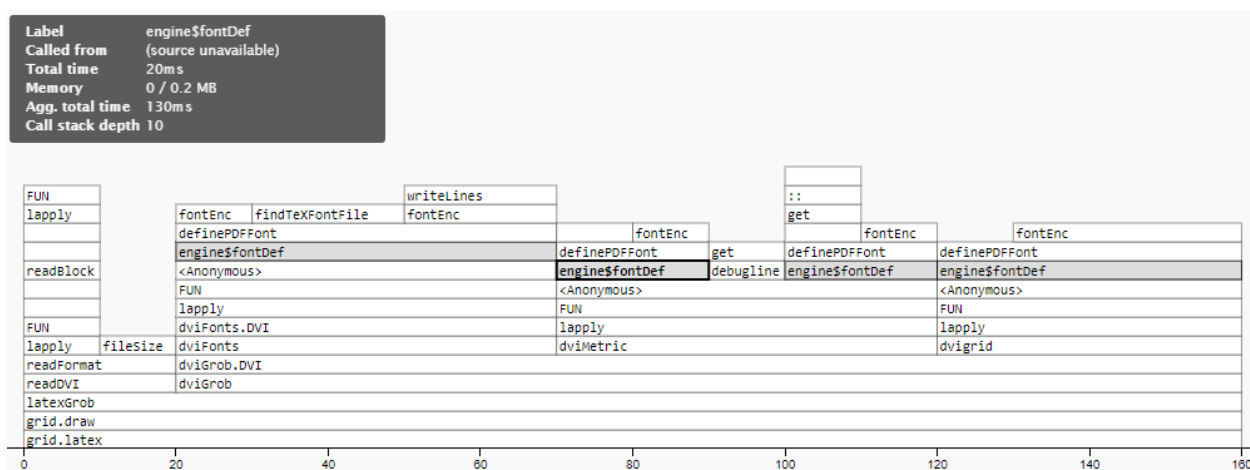


Figure 7: Screenshot of `profvis::profvis()` output for the code `grid.latex("$x - \\mu$")` in `dvir` version 0.2-1, highlighting `engine$fontDef`.

The effect of this is very obvious in figure @ref(fig:profilingSimpleProfvis_0.2-1_highlight) which is the same as figure @ref(fig:profilingSimpleProfvis_0.2-1) but highlights the time spent in `engine$fontDef`. The wrappers for the font, metric and grid sweeps are `dviFonts()`, `dviMetric()` and `dvigrid()` respectively (sixth call from the bottom of the stack). Here we can see nearly all of the time spent in the metric and grid sweeps are actually redoing the font sweep.

The change to be made was simply stopping the metric and grid sweeps from doing the font sweep again.

The font sweep looks in the DVI file for op codes 243 to 246. These are the op codes for font definitions and define the name of a font and give it an identifier to reference in the DVI file when it wants to use that font to display a character.

The following code is taken from the `dvir` package, showing where the metric and grid sweeps also redid

the font sweep. `op_font_def` is a function which takes the font definition in the DVI file related to that instance of the op code and searches for and records the font information.

```
metric_info_243 <- op_font_def
grid_op_243 <- op_font_def
```

The following code shows what the code was changed to in `dvir` version 0.2-2. `op_ignore` is an empty function, so when the metric or grid sweeps comes across that op code, they now do nothing.

```
metric_info_243 <- op_ignore
grid_op_243 <- op_ignore
```

Unfortunately these changes by themselves caused an error when running `grid.latex()`. This is because one task undertaken before the font sweep is to reset or overwrite the global fonts list (which the font sweep then writes to). The metric and grid sweeps were also doing this even though it was only intended for it to be done by the font sweep. This meant after the font sweep was completed it was overwritten by the metric and grid sweeps and so when `dvir` tried to draw the characters there was no font information to refer to.

The resetting of the global fonts list was initiated when the sweeps passed op code 247 in the DVI file, which is the preamble at the start of every DVI file. Setting the metric and grid sweeps to do nothing when they pass the preamble of the DVI file, again by way of `op_ignore`, solved this problem as the global fonts list created by the font sweep is now not overwritten.

```
metric_info_247 <- op_ignore
grid_op_247 <- op_ignore
```

To quantify the impact this has on code speed we recorded the time to run our examples 20 times, after an initial run to compile the package after it was loaded.

```
## Warning: `expand_scale()` is deprecated; use `expansion()` instead.
```

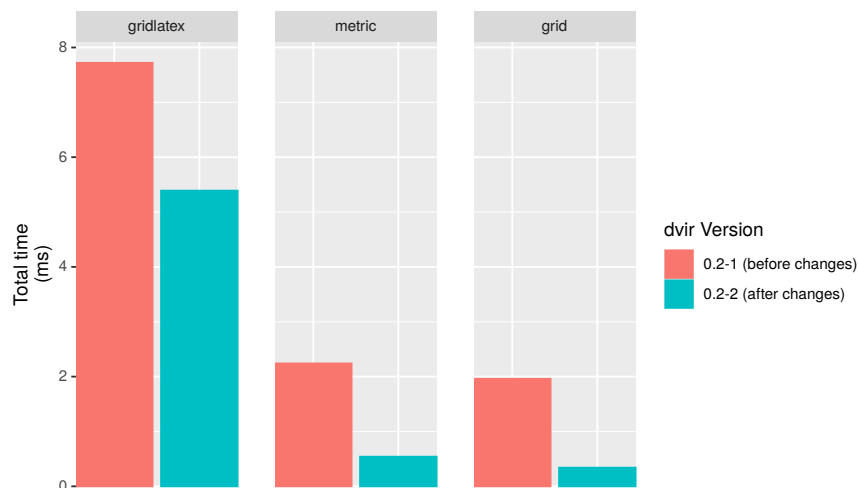


Figure 8: The total time spent in the `'grid.latex()'` function, metric sweep and grid sweep before and after these changes, over 20 runs of our simple example.

For the simple example the time spent in the metric and grid sweeps have decreased by 76% and 83%

respectively. The speed of the overall `grid.latex()` call has decreased by 30%.

```
## Warning: `expand_scale()` is deprecated; use `expansion()` instead.
```




Figure 9: The total time spent in the ‘`grid.latex()`’ function, metric sweep and grid sweep before and after these changes, over 20 runs of our complicated example.

Similarly for the more complicated example in figure @ref{fig:speedUp1Yee} the metric and grid sweeps have decreased by 83% and 89% respectively, with `grid.latex()` overall taking 47% less time.

5 Code speed (part 2) - font caching

In the previous section, the speed of `dvir` was increased by stopping it doing something “silly”. The profiling results from the previous section demonstrate how much time `grid.latex()` spends dealing with fonts so this section is about making `dvir` “smarter” by caching the information it gathers about fonts between `grid.latex()` calls.

For example, creating figure ?? requires three `grid.latex()` calls for the same symbol ω . Each time, the following font definition appears in the DVI file for `grid.latex()` to search for during its font sweep.

```
fnt_def_1    fontnum=17, checksum=-1209964637, scale=786432, design=786432,
             fontname=cmmi12
```

Caching fonts requires several steps:

- During the font sweep, search all font definitions as per usual.
- For each font definition, check if the font has been cached already.
 - If yes, do nothing. Refer to the existing font information in the cache when the needs to be used
 - If no, search for and store the font information as per usual.

The font information is stored in R as a list. Every time `grid.latex()` processed a “preamble” op code in a DVI file, that is, every time `grid.latex()` was called, the font list was initialised as an empty list in the `dvir` namespace with `dvir::set("fonts", vector("list", 255))`. To stop this initialisation happening by default, the following code was put in its place.

```
if (dvir::get("initFonts") || is.null(dvir::get("fonts"))) dvir::set("fonts", vector("list", 255))
```

Note in this new statement the logical variable `initFonts`, stored in the `dvir` namespace. This lets a user choose whether to initialise the font cache on a `grid.latex()` call. If `initFonts` is `FALSE` or the font list does not exist, for example after the package is loaded, the font list is initialised. Otherwise, the existing font cache remains. `initFonts` was added as an argument to `grid.latex()` with its default value being equal to the global option `dvir.initFonts`.

In a DVI file, each font definition gives the font a identification number, for example `fontnum=17`. In all the DVI files examined as part of this project, this font identification number is the same when that particular font is used again in another `grid.latex()` call. This number is used as the index of the fonts list where its information is stored.

When the font sweep passes a font definition in the DVI file it checks first whether the index of the font list matching the font number of the new font contains any existing font information:

- If no, then the font information is written to the font list as usual.
- If yes, then it checks if the existing font information is the same as the font information in the DVI file.
 - If yes, then it does nothing.
 - If no, then it overwrites the fonts as per usual.

The result of this is that across multiple `grid.latex()` calls requiring the same fonts, any future uses of that font are much faster after the font has been used once. Importantly it also has a “fail safe” where if there happens to be a case where the font number was used with a different font in a previous `grid.latex()` call, the font cache is overwritten with the new font so the correct font will be used.

The updated font list is stored in the `dvir` environment to be retrieved the next `grid.latex()` finds a font definition during a font sweep.

The only weakness of this process is if in a *single* DVI file different fonts are defined with the same font number. In this case the second font will be stored in the font cache in that index and so will be used in place of the first font. No instances of this occurring were discovered during this project.

- **[Example:** Do example counting how many times same font definition is repeated in thomas yee’s example]

To examine the effect of these changes on code speed we performed profiling as was detailed in the previous section.

```
## Warning: `expand_scale()` is deprecated; use `expansion()` instead.
```

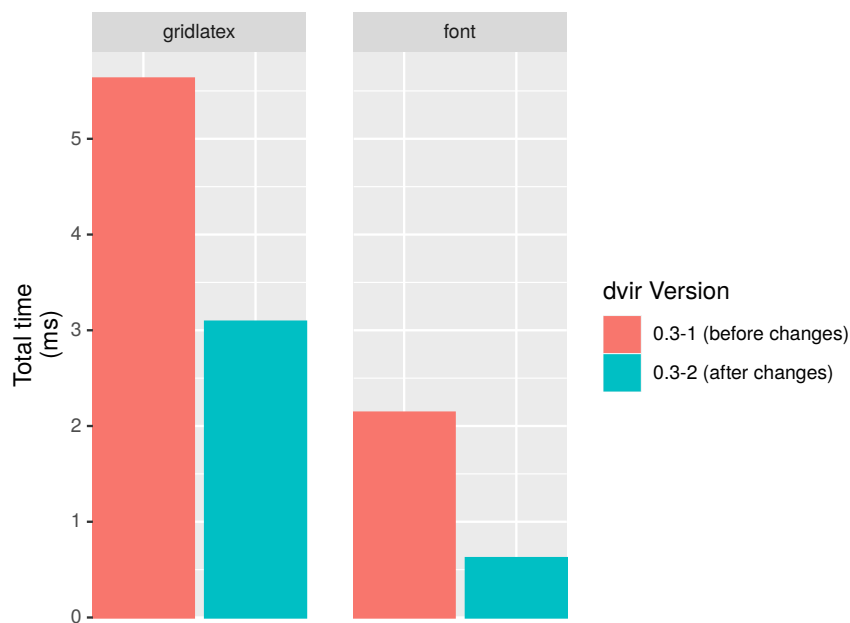


Figure 10: The total time spent in the ‘grid.latex()’ function and font sweep

```
## Warning: `expand_scale()` is deprecated; use `expansion()` instead.
```

5.1 Profiling environment specifications

The exact results obtained in this and the previous section are specific to the computing environment used. Specific details are provided below. The sampling nature of profiling (intermittent recording of the call stack) will give different results every time it is done.




Figure 11: The total time spent in the ‘grid.latex()’ function and font sweep

The profiling results are dependent on the computer setup used and could change depending on the exact computing environment in which the `dvir` package is used.

The profiling results in this report, in this and the previous section, were calculated with the following setup:

- A virtual machine via Oracle VM Virtualbox
- Virtual machine running Ubuntu 18.04.5 LTS
- R version 3.4.4
- `dvir` package versions as described with the profiling results

6 Linear gradient fills

6.1 TikZ and dvir

TikZ is a TeX package that allows drawing of pictures and diagrams in TeX documents [reference TikZ report/description]:

- [Example: simple TikZ drawing (circles with labels, and an arrow maybe)]
- [Example: more complicated TikZ drawing, maybe with colouring and stuff]

The original DVI specification only needed to account for text and typesetting (and the most basic of rectangles) and so was not designed with drawing and graphics in mind. The type of instruction in the DVI file are labelled with an “op code”. Each op code described a type of instruction like defining fonts, setting characters to display and vertical and horizontal cursor movements. There were four op codes however, called *DVI specials*, that can contain almost any form of instruction or values needed, such as text colour, to create a document based on the DVI file, such as Postscript or PDF.

The TikZ package uses these DVI specials to describe shapes, drawings and colours in PGF (portable graphics format) which can be translated to instructions for other viewing formats, like Postscript, PDF or SVG. How the instructions are translated is controlled by a TikZ driver. The `dvir` package includes its own TikZ driver to translate the drawing instructions into a form useful to draw the things with R grid graphics [reference Paul dvir TikZ report].

Some TikZ features were not implemented though, notably the ability to have fill colours of shapes as linear or radial gradients or patterns. The primary reason for this is that R did not support these types of fills but the latest R release in May 2021, version 4.1.0, provides support for these fills in the `grid` package, on which `dvir` is built.

- [Example: replicate one of the above examples in R]
- [Example: TikZ radial gradient fill example]
- [Example: Make same TikZ example as above in R with dvir (obviously fill will be blank)]
- [Example: Use R 4.1.0 to make a linear gradient in a shape]

As it is, the TikZ driver simply ignores any gradient or pattern fill information when creating the DVI file for `dvir`.

- [Example: Use `grid.tikzpicture()` for picture with gradient fill in text, but resulting R graphic does not have fill]

6.2 Implementing TikZ linear gradient fills in dvir

The following steps are required to implement these TikZ fills in `dvir`:

1. Add the fill information (like the gradient colours and their locations) to the DVI file created by `dvir`
2. Store this fill information during a parse by `dvir` to read the DVI file


3. Add the fill information when drawing the shape in R

To tackle step 1, we need to update the `dvir` TikZ driver file to include information about the gradient and pattern fills. As the `dvir` TikZ driver file is based on the SVG TikZ driver file, the SVG support for TikZ fills was used as a base to edit to make it specific to `dvir`.

The information we require for the gradient fills from TikZ via the DVI file is as per the arguments for the `grid::linearGradient(...)`, which is used as an argument to `grid::gpar(fill = linearGradient(...))`, which itself is an argument to a `grid` drawing function, for example `grid::grid.rect(..., gp = gpar(fill = linearGradient(...)))`. The most important parts of defining a linear gradient fill is the colours and stops of the gradient fill. The stops of a gradient fill are the locations along the length of a gradient fill where the specified colours are. In between the stops, the gradient between stop colours either side occurs.

The `colours` and `stops` arguments of `linearGradient()` are simply vectors of colours (a character vector of colour names or hexadecimal RGB values) and locations of those colours as a proportion of the distance between the start and end points of the gradient respectively. This obviously guides us as to what information we need to get from TikZ in the DVI file so we can pass it to `dvir`.

Let us consider a simple example, a rectangle with an orange to green linear gradient fill:



The following is an extract of the DVI file when the rectangle above is generated using the SVG DVI driver included with the common T_EX distributions, `pgfsys-dvisvgm.def`. It has been edited slightly for readability by adding line breaks.

```
xxx1      k=67
          x=dvisvgm:raw <g transform="matrix(1,0,0,1,56.90549,28.45274)">{?nl}
xxx1      k=67
          x=dvisvgm:raw <g transform="matrix(2.26802,0,0,1.134,0.0,0.0)">{?nl}
xxx1      k=66
          x=dvisvgm:raw <g transform="matrix(0.0,1.0,-1.0,0.0,0.0,0.0)">{?nl}
xxx4      k=425
          x=dvisvgm:raw <linearGradient id="pgfsh2" gradientTransform="rotate(90)">{?nl}
                                <stop offset=" 0.0" stop-color=" rgb(0.0%,100.0%,0.0%) ">{?nl}
                                <stop offset=" 0.25" stop-color=" rgb(0.0%,100.0%,0.0%) ">{?nl}
                                <stop offset=" 0.5" stop-color=" rgb(50.0%,75.0%,0.0%) ">{?nl}
                                <stop offset=" 0.75" stop-color=" rgb(100.0%,50.0%,0.0%) ">{?nl}
                                <stop offset=" 1.0" stop-color=" rgb(100.0%,50.0%,0.0%) ">{?nl}
                                </linearGradient>{?nl}
xxx1      k=57
          x=dvisvgm:raw <g transform="translate(-50.1875,-50.1875)">
xxx1      k=97
```

```
x=dvisvgm:raw <rect width="100.375" height="100.375"
style="fill:url(#pgfsh2); stroke:none"/>{?nl}
```

We can see from this that the linear gradient definition with stops and colours is defined within a `<linearGradient>` element and given an `id` attribute. In the `<rect>` element a CSS style definition sets the fill of the rectangle by referring to the `id` of the previously defined definition. In the linear gradient definition there are colours defined as RGB values and their respective stops so now we need to get the `dvir` driver file to extract the same information in a “R-friendly” form. There are several \TeX macros that were defined to do this. These are based on the `pgfsys-common-svg.def` and `pgfsys-dvisgm.def` SVG drivers that come with most \TeX distributions.

A “wrapper” macro of what to do when a gradient fill is requested. Within this, the definition of the fill is created and sent to the DVI file, followed by a rectangle with a fill specified by that definition. For clarity, a line stating when the gradient fill is defined and then again when it is used has been added to the DVI file but these can be removed.

```
\def\pgfsys@shadinginsidepgfpicture#1{%
  #1%
  \pgfsysprotocol@literal{SHADING BEING DEFINED: ShadDefID = \the\pgf@sys@dvir@objectcount}%
  \pgf@sys@dvir@sh@defs%
  \pgf@process{\pgf@sys@dvir@pos}%
  \pgf@xa=-.5\pgf@x%
  \pgf@ya=-.5\pgf@y%
  \pgfsysprotocol@literal{<g transform="translate(\pgf@sys@tonumber{\pgf@xa},\pgf@sys@tonumber{\pgf@ya})"}
  \pgfsysprotocol@literal{SHADING BEING USED: ShadDefID = \the\pgf@sys@dvir@objectcount}
  \pgf@sys@dvir@sh%
}
```

In basic cases, a horizontal linear gradient is defined as a vertical gradient with a 90 degree rotation which is why this is called “vert” shading, as it is in the SCG driver file. The definition and use of the gradient specified above are collated here. The stop positions (proportion along the length of the gradient for which a colour specified) and their respective colours are collated and the literal text to build the gradient definition and rectangle using that definition. The biggest difference from the SVG driver was that R needs a vector of stops and a separate vector of colours as arguments to `linearGradient()`. SVG specifies the position and colour of each stop together.

```
\def\pgfsys@vertshading#1#2#3{%
  {%
    \pgf@parsefunc{#3}%
    \global\advance\pgf@sys@dvir@objectcount by1\relax%
    \pgf@sys@dvir@shading@stop%
    \pgf@sys@dvir@shading@stopcolours%
    \expandafter\xdef\csname @pgfshading#1!\endcsname{%
      \def\noexpand\pgf@sys@dvir@sh@defs{\noexpand\pgfsysprotocol@literal{\pgf@sys@dvir@thestops}}%
      \def\noexpand\pgf@sys@dvir@sh{\noexpand\pgfsysprotocol@literal{<rect
        width="\pgf@sys@tonumber{\pgf@y}"
```

```

        height="\pgf@sys@tonumber{\pgf@x}"
        style="fill:url(\noexpand\#pgfsh\the\pgf@sys@dvir@objectcount);
        stroke:none"/>\noexpand\pgf@sys@dvir@newline}}%
\def\noexpand\pgf@sys@dvir@pos{\noexpand\pgfpoint{\the\pgf@y}{\the\pgf@x}}%
}%
}%
}

```

This macro allows us to iteratively add stop positions or colours to the ones we have already gathered.

```

\let\pgf@sys@dvir@thestops=\pgfutil@empty
\def\pgf@sys@dvir@addtostops#1{%
  \edef\pgf@temp{#1}%
  \expandafter\expandafter\expandafter\def
  \expandafter\expandafter\expandafter\pgf@sys@dvir@thestops
  \expandafter\expandafter\expandafter{\expandafter\pgf@sys@dvir@thestops\expandafter\space\pgf@temp}%
}

```

The following macros process all the stop locations, collating them into an R friendly vector suitable to parse to `linearGradient()`

```

\def\pgf@sys@dvir@shading@stop{%
  % Step 1: Compute 1/\pgf@sys@shading@end@pos
  \pgf@x=\pgf@sys@shading@end@pos\relax%
  \c@pgf@counta=\pgf@x\relax%
  \divide\c@pgf@counta by4096\relax%
  % Step 2: Insert stops locations
  \pgf@sys@dvir@addtostops{stops={}%
  \expandafter\pgf@sys@dvir@shading@dostoplocations\pgf@sys@shading@ranges%
  % dummy for end:
  {\pgf@sys@shading@end@pos}{\pgf@sys@shading@end@pos}}%
  \pgf@sys@dvir@addtostops{)}%
}

```

```

\def\pgf@sys@dvir@shading@dostoplocations#1{%
  \edef\pgf@test{#1}%
  \ifx\pgf@test\pgfutil@empty%
  \else%
    \expandafter\pgf@sys@dvir@shading@dostoplocation\pgf@test%
    \expandafter\pgf@sys@dvir@shading@dostoplocations
  \fi%
}

```

```

\def\pgf@sys@dvir@shading@dostoplocation#1#2#3#4{%
  % #1 start pos
  % #2 end pos

```

```

% #3 start rgb
% #4 end rgb
\pgf@x=#1%
\pgf@x=16\pgf@x%
\divide\pgf@x by \c@pgf@counta\relax%
\expandafter\pgf@sys@dvir@addtostops{\pgf@sys@tonumber\pgf@x}%
}

```

Similar to the above, these collate all the colours of the stops into an R friendly vector.

```

\def\pgf@sys@dvir@shading@stopcolours{%
% Step 1: Compute 1/\pgf@sys@shading@end@pos
\pgf@x=\pgf@sys@shading@end@pos\relax%
\c@pgf@counta=\pgf@x\relax%
\divide\c@pgf@counta by4096\relax%
% Step 2: Insert stops RGB colours
\pgf@sys@dvir@addtostops{, colours={}%
\expandafter\pgf@sys@dvir@shading@dostopcolours\pgf@sys@shading@ranges%
% dummy for end:
{\pgf@sys@shading@end@rgb}{\pgf@sys@shading@end@rgb}{\pgf@sys@shading@end@rgb}}%
\pgf@sys@dvir@addtostops{)}%
}

```

```

\def\pgf@sys@dvir@shading@dostopcolours#1{%
\edef\pgf@test{#1}%
\ifx\pgf@test\pgfutil@empty%
\else%
\expandafter\pgf@sys@dvir@shading@dostopcolour\pgf@test%
\expandafter\pgf@sys@dvir@shading@dostopcolours%
\fi%
}

```

```

\def\pgf@sys@dvir@shading@dostopcolour#1#2#3#4{%
% #1 start pos
% #2 end pos
% #3 start rgb
% #4 end rgb
\expandafter\pgf@sys@dvir@shading@dorgb#3%
}

```

This is a helper function to return formatted RGB values.

```

\def\pgf@sys@dvir@shading@dorgb#1#2#3{%
\pgf@sys@dvir@color@rgb#1,#2,#3\relax%
\pgf@sys@dvir@addtostops{\pgf@sys@dvir@prepared}%
}

```

A counter is defined in order to give each gradient definition unique identifier.

```
\newcount\pgf@sys@dvir@objectcount
```

The result of these changes mean that the linear gradient definition is specified in the DVI file in a form usable by R, sans some commas:

```
xxx1      k=118
          x=dvir:: stops=( 0.0 0.25 0.5 0.75 1.0 ) ,
                      colours=( rgb(0,1,0) rgb(0,1,0) rgb(0.5,0.75,0) rgb(1,0.5,0) rgb(1,0.5,0) )
```

Compare this with the linear gradient definition from earlier:

```
xxx4      k=425
          x=dvisvgm:raw <linearGradient id="pgfsh2" gradientTransform="rotate(90)">{?nl}
                      <stop offset=" 0.0" stop-color=" rgb(0.0%,100.0%,0.0%) "/>{?nl}
                      <stop offset=" 0.25" stop-color=" rgb(0.0%,100.0%,0.0%) "/>{?nl}
                      <stop offset=" 0.5" stop-color=" rgb(50.0%,75.0%,0.0%) "/>{?nl}
                      <stop offset=" 0.75" stop-color=" rgb(100.0%,50.0%,0.0%) "/>{?nl}
                      <stop offset=" 1.0" stop-color=" rgb(100.0%,50.0%,0.0%) "/>{?nl}
                      </linearGradient>{?nl}
```

Unfortunately this was the extent this project could explore implementing TikZ gradient fills in dvir. You may notice in both the SVG and dvir DVI output that the colour at stops 0 and 0.25 are the same, as is the colour at stops 0.75 and 1. This is related to a larger issue of TikZ using a number of transformations on the “simple” linear gradient definition above which is then clipped to the shape the fill is for. These transformations would require some very detailed work to translate them into a set of instructions to perform the same in R before they are used in `linearGradient()`.

An additional complication is that TikZ will create shapes with these fills by first drawing the shape with a transparent fill and then drawing another shape with the gradient fill clipped to the size of the first shape. R however requires the the fill information as an argument when drawing the shape. In R, any transformations of the gradient, as TikZ does, will need to be done *before* it is parsed to the drawing function.

Radial gradient fills have similar problems with how they are manipulated by TikZ, however the work done so far on linear gradients could be easily replicated for them.

This work is still within step 1 of the workflow in section 6.2. To continue with the implementation, aside from the above obstacles, dvir would need to do another sweep of the DVI file, like the font sweep, to achieve step 2. After any transformations or manipulations are made, the gradient definition can be called by a unique identifier to parse the stored fill definitions to the drawing function used as in step 3.

7 Text baselines

7.1 The problem

Text characters have a baseline, that is, a horizontal line on which the characters naturally sit so all the letters appear to be in line with each other. Some letters, like a lower case “p” or “j”, have a “descender”. A descender is the part of a character that sits *below* the baseline.

- [Example - showing a letter, with the baseline marked and also showing what part is the “descender” part]

`grid.text()` accounts for this baseline so when you bottom align text at a certain y-value, the descenders will actually fall below the y-value we defined, despite the bottom justification. Unfortunately, `dvir` does not account for text baselines and will do any alignment with the bounding box of the text.

- [Example - combining `grid.text()` and `grid.latex()` on same y-value to show the difference between them]
- [Example - more complicated example, say with several pieces of multiline text (only if I can demonstrate later that my function works for aligning them!)]

To fix this, `dvir` needs to obtain or calculate a value for the baseline for any given piece of text to offset the bounding box when it is drawing text. All of `dvir`’s information comes from the DVI file we either need to extract a baseline value from the DVI file itself, or calculate it *from* information in the DVI file. Unfortunately plain DVI files do not state a baseline value, rather they only detail specific placement of each character, so we will explore some possible methods to obtain a baseline value. These methods are referred to as algorithms from here on due to their heuristic nature.

7.2 Implementation

To explore the algorithms detailed below and evaluate their usefulness an R function, `baselines()` has been created. This function takes several arguments, including the baseline selection algorithm as detailed below, any additional information needed for that particular algorithm, and the `TeX` code as you would use with `grid.latex()`. The output of this function is the distance, or in some case distances, from the bottom of the bounding box of the text to the possible baseline value. These distances are returned as `grid` units.

Once the baselines have been calculated the bounding box of the text can be bottom aligned with the y-value specified but then moved down by the value of the baseline. This means the baseline of the text will be at the specified y-value.

This function has been written to easily allow integration of other algorithms and most of the function could be directly implemented in the `dvir` package, should this baseline algorithm feature be implemented into `dvir` formally.

7.3 Our potential solutions

Several different algorithms were explored to calculate the baseline for different types of text that could be used with `grid.latex()`. These algorithms are detailed here.

7.3.1 alex algorithm

This is a simple algorithm which was determined after inspection of some DVI files. In every DVI file, there is a statement specifying the size of the bounding box of the text.

- [Example of the HiResBoundingBox statement here]

After this statement there appears to consistently be a downward move equal to the height of the bounding box of text (from the top left of the bounding box to the bottom left), and then a move upward before the first character is drawn. This algorithm takes the cursor location after that upward move to be the baseline. In instances where the entire text has no descenders (i.e. the baseline is the bottom of the bounding box) there will be no upward move before the first character and a value of 0 is returned as the baseline.

```
baselineValue <- baselines(tex = "$x - \\mu$", algorithm = "alex")
baselineValue
```

```
## [1] 127431scaledpts
```

For plain text this algorithm works quite well. Unfortunately for many equations, particularly ones with superscripts or subscripts, and multiline text the first upward move is *not* to where the baseline is.

test

testing

var

varying

$\sum_{n=1}^{\infty} 2^{-n} = 1$

$\sum_{n=1}^{\infty} 2^{-n} = 1$

$\sum_{n=1} 2^{-n} = 1$

The equation is $x + \frac{\mu^2}{2}$

The equation is $x + \frac{\mu}{2}$

$\frac{\mu^2}{2} + x$ is the equation

$x + \frac{\mu}{2}$ is the equation

Paragraph, with
some line wrap-
ping!

Figure 12: Using the ‘alex’ method to calculate the baseline.

7.3.2 dviMoves algorithm

This is an extension of the `alex` algorithm. Rather than only taking the location after the second ‘down’ move, this algorithm keeps track of *all* the up and down moves of the cursor. The motivation behind this is that the upward and downward “moves” in the DVI file reflect the cursor moving to the baseline value of the next character to be typeset. Once again we assume that the first downward move after the “HiResBoundingBox” statement is from the top to the bottom of the bounding box and so this algorithm only returns the upward and downward cursor moves from there, however as DVI files have the ability to save the current cursor location, move around a bit, then reset back to the saved location, all up and down moves are recorded from the start of the DVI file.

There are two complications with this method:

- As it returns *all* the vertical positions the cursor moves to there are many possible baseline values returned. Any more than one means we have to decide which of them to choose.
- There are often several up and downward moves in the DVI file between typeset characters so in between the “useful” baselines there can be some which are not so useful or duplicates.

To account for these considerations, along with the `dviMoves` algorithm, the function also allows a choice of method to select a *single* baseline out of the many returned by the algorithm. These methods are:

7.3.2.1 dviMoves selection method all

This selection method will return all the baseline values as determined by the algorithm. This is useful for drawing the text with all the baseline values calculated from this method.

7.3.2.2 dviMoves selection method index

When this selection method is chosen, another argument to the function, `dviMovesSelection`, is used to specify a numeric index. The baseline value corresponding to that index is returned.

7.3.2.3 dviMoves selection method bottomUp

This is similar to the `index` method but the baseline values are first ordered from smallest to largest, before using the `dviMovesSelection` argument to select the index of the baseline value to return.

7.3.2.4 dviMoves selection method nextChars

This method will return the first baseline value just after a specified character in the argument `dviMovesSelection`. Characters are specified by a number, entered as a character string. This method is limited for several reasons:

- It cannot return a baseline value for any instance of a character after the first instance
- If multiple cursor moves are made before a character is drawn, the earlier moves are still recorded but with an empty character label
- Some characters are hard to type on a “normal” keyboard, for example μ

- The DVI file does not actually contain the character, but rather the numeric index of it in the particular font. This is why a number has to be used to select which character it is.




Figure 13: Using the ‘dviMoves’ method to calculate the baseline.

`cat(latexCode)`

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

Paragraph, with
some line wrap-
ping!

7.3.2.5 Future options

Some other selection methods could be:

- Taking the baseline value closest to the baseline value determined by one of these other algorithms
- `prevChar`
- extend `dviMoves` to only record position just before a character is typeset?

7.3.3 preview algorithm

This algorithm uses the `preview` package in \LaTeX . [Reference TuG report here] When this algorithm is selected, the following code is added to the header of the DVI file:

```
\usepackage[active,showbox,tightpage]{preview}
\def\showbox#1%%
{\immediate\write16{MatplotlibBox:(\the\ht#1+\the\dp#1)x\the\wd#1}}
```

This results in the following line being printed in the log file.

```
MatplotlibBox:(8.04175pt+3.00005pt)x62.81718pt
```

The baseline value in this case is the value just after the `+` and can be obtained by searching the log file for this line starting with `MatplotlibBox:` and returning the value just after the `+`.

7.3.4 dvipng algorithm

The `dvipng` program is used to turn DVI files into PNG images. An option of `dvipng` is `--depth` which returns the baseline value in pixels. [reference dvipng package - maybe where you get it from linux?]

The system call `dvipng --depth test.dvi` returns `depth=4` within its output so after recording the output, a simple search for `depth=` will find the baseline value.

7.3.5 Other potential algorithms

Some other potential algorithms but have not yet been implemented in this function include:

- “Just ask \TeX ”

test

testing

var

varying

$$\sum_{n=1}^{\infty} 2^{-n} = 1$$

$$\sum_{n=1}^{\infty} 2^{-n} = 1$$

$$\sum_{n=1} 2^{-n} = 1$$

The equation is $x + \frac{\mu^2}{2}$

The equation is $x + \frac{\mu}{2}$

$\frac{\mu^2}{2} + x$ is the equation

$x + \frac{\mu}{2}$ is the equation

Paragraph, with
some line wrap-
ping!

Figure 14: Using the ‘preview’ method to calculate the baseline.

test

testing

var

varying

$$\sum_{n=1}^{\infty} 2^{-n} = 1$$

$$\sum_{n=1}^{\infty} 2^{-n} = 1$$

$$\sum_{n=1} 2^{-n} = 1$$

The equation is $x + \frac{\mu^2}{2}$

The equation is $x + \frac{\mu}{2}$

$\frac{\mu^2}{2} + x$ is the equation

$x + \frac{\mu}{2}$ is the equation

Paragraph, with
some line wrap-
ping!

Figure 15: Using the ‘dvipng’ method to calculate the baseline.

7.4 Discussion of algorithms

Of the algorithms considered, the `dviMoves` algorithm performed best. While most of these algorithms perform well for some or most of the examples, the `dviMoves` algorithm performs well for *all* the examples, notably for giving the option to align the baseline of *any* line of multi-line text. Utilising the fact it returns all baseline values for all characters, it is possible to align, for example, with any character in a mathematical equation, whether it be of a different size or a superscript or subscript.

- [Example: Show examples from start of this section with baselines being used to make them aligned!]

The `preview` algorithm also calculated a baseline well for all examples, except the multi line text. An implementation of this in `dvir` could perhaps use both the `dviMoves` and `preview` algorithms to calculate the baselines, with the default selection method being the baseline value from `dviMoves` that is closest to the baseline value calculated by the `preview` algorithm.

8 Conclusion

\TeX and its extensions have more functionality than could realistically be added to R but by targeting the most useful features of \TeX , as `dvir` did originally and as has been continued in this project, many use cases are accommodated. The changes discussed in this report have made the `dvir` package better especially for displaying text and equations, graphics with many `grid.latex()` calls, or even when developing graphics with `grid.latex()` which requires rerunning code with small tweaks to get the perfect graphic.

9 References

Murrell, Paul. 2018. “Revisiting Mathematical Equations in R: The ‘dvir’ Package.” 2018-08. Department of Statistics, The University of Auckland.

10 Apendices

Associated files related to this project can be found on GitHub on the link in the Executive Summary.